Sampling and Inference for Beta Neutral-to-the-Left Models of Sparse Networks

Benjamin Bloem-Reddy, Adam Foster, Emile Mathieu, Yee Whye Teh

Department of Statistics, University of Oxford

Contents

Temporal networks

Examples

▶ Messages sent between people over time

▶ Protein-protein interactions

Temporal networks

Temporal networks

Temporal networks

Temporal networks

Temporal networks

Edges and vertices

Edges and vertices

Edges and vertices

Edges and vertices

Sparsity

Power law degree distribution

$p(d) \propto d^{-\eta}$, for $\eta > 1$

$\eta = 1.5$

$\eta = 2.5$

$\log p(d)$

10^0 10^1 10^2 10^3 10^4

Log-degree, $\log d$

Sparsity and power law

$$\textbf{Sublinear} \text{ sparsity} \quad \Longleftrightarrow \quad \eta \in (1,2)$$
$$\textbf{Linear} \text{ sparsity} \quad \Longleftrightarrow \quad \eta > 2$$

Empirical study

SNAP dataset [1]	# of vertices	# of edges
Ask Ubuntu	159,316	964,437
UCI social network	1,899	20,296
⋮	⋮	⋮

Ask Ubuntu

UCI social network

Generative models for network data
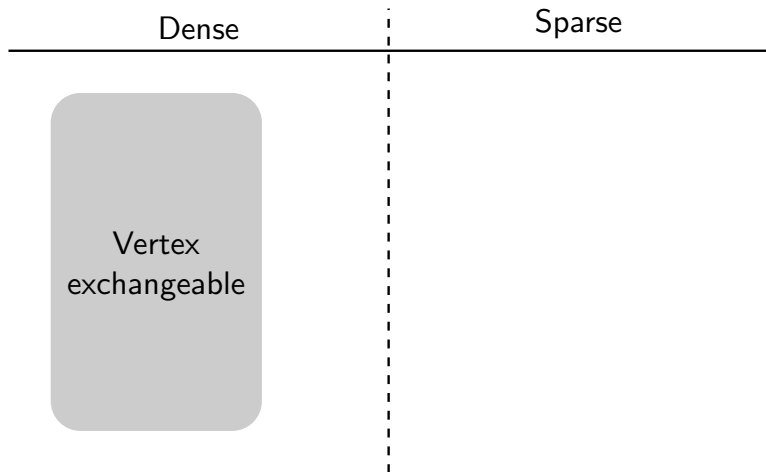
Dense	Sparse
Vertex exchangeable	

Generative models for network data

Generative models for network data

Generative models for network data
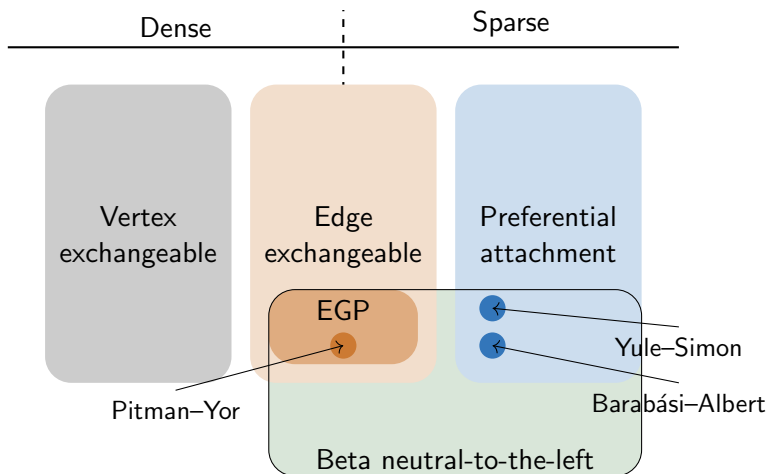
Edge exchangeable models [2], [3]

Factorizable ("rank one") paintbox representation

Sequence of edges: $((3,6),(1,5),(1,7))$

Factorizable ("rank one") paintbox representation

Sequence of edges: $((3, 6), (1, 5), (1, 7))$

Paintbox representation

Consequences for edge exchangeable models:

▶ Rate of vertex arrival gets slower and slower
$\Rightarrow$ sublinear sparsity: #vertices $= o(\text{#edges})$

▶ Edges pile up
$\Rightarrow$ linear scaling of degrees: $d_{j,n} = \Theta(n)$

Models

dense   Sparse
         sublinear   linear

Vertex exchangeable

Edge exchangeable

Pitman Yor

Preferential attachment
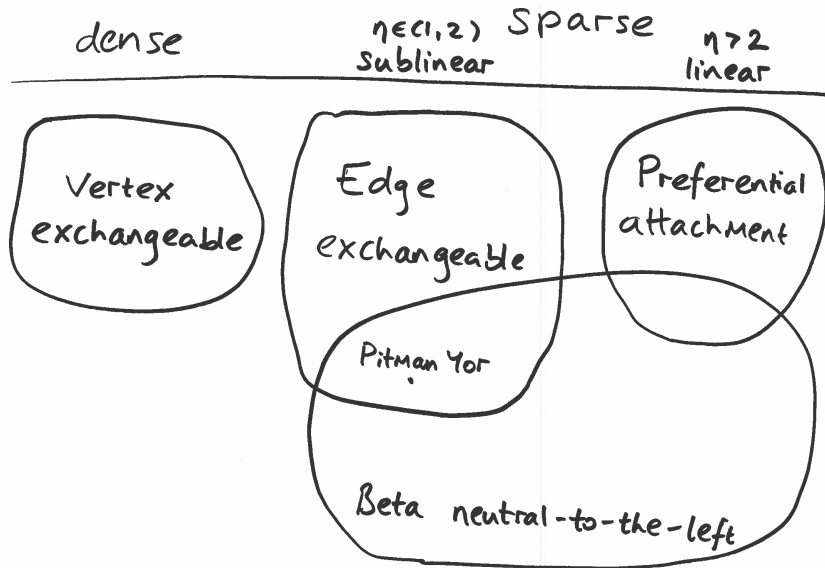
Models

Models

Beta Neutral-to-the-left Model [4]

1. Generate arrival times $1 = T_1 < T_2 < T_3 < \ldots$ in any way.
2. Generate ends of edges sequentially:

G_n

Is the next arrival time equal to $n + 1$?

yes ← → no

G_{n+1}

New vertex

$\mathbb{P}[\to j] \propto \deg_{j,n} - \alpha$

G_{n+1}

Sequence of paintboxes representation

Sampling and inference

Why a paper on sampling and inference for BNTL models?

Sampling and inference

Why a paper on sampling and inference for BNTL models?

- ▶ Inference is notoriously difficult for *non-exchangeable* structures
- ▶ Need to identify *exchangeable substructures*

Exchangeable substructure

Gibbs structure

The joint probability has **Gibbs structure** due to left-neutrality

$$P(\text{graph}|\mathbf{T}) = \prod_{j=1}^{K} P(\text{choose } j \ d_j - 1 \text{ times out of } n - T_j \text{ trials })$$

- $K = \#\text{vertices}$
- $n = \#\text{edges}$
- $d_j = \text{degree of vertex } j$
- $T_j = \text{arrival time of vertex } j$

Gibbs structure

Explicitly,

$$p(\text{graph}|\mathbf{T}) = \frac{\Gamma(d_1 - \alpha)}{\Gamma(n - K\alpha)} \prod_{j=2}^{K} \frac{\Gamma(T_j - j\alpha)\Gamma(d_j - \alpha)}{\Gamma(T_j - 1 - (j-1)\alpha)\Gamma(1 - \alpha)}$$

- $K = \#\text{vertices}$
- $n = \#\text{edges}$
- $d_j = \text{degree of vertex } j$
- $T_j = \text{arrival time of vertex } j$

Available data

Observation	Unobserved variables
Entire history	$\alpha, \phi, \boldsymbol{\Psi}$
Degrees in arrival order	$\alpha, \phi, \boldsymbol{\Psi}, \mathbf{T}$
Snapshot	$\alpha, \phi, \boldsymbol{\Psi}, \mathbf{T}, \sigma$

- α = BTNL parameter $\in (-\infty, 1)$
- ϕ = arrival distribution parameters
- $\boldsymbol{\Psi}$ = latent sociabilities
- $\mathbf{T}$ = arrival times
- σ = arrival order

Gibbs sampler

Variable	Gibbs sampling scheme
α	MCMC, e.g. slice sampling
ϕ	Depends on arrival dist. family Λ_ϕ
$\boldsymbol{\Psi}$	$\Psi_j|\text{graph}, \boldsymbol{\Psi}_{\setminus j} \sim \text{Beta}(d_j - \alpha, \bar{d}_{j-1} - (j-1)\alpha)$ where $\bar{d}_j = \sum_{i=1}^{j} d_i$ can marginalise out $\boldsymbol{\Psi}$
$\boldsymbol{T}$	Simple update for T_j, can't move past neighbours
σ	Initialise in descending degree order use M-H with adjacent swap proposal $\sigma_j \leftrightarrow \sigma_{j+1}$ fast to compute due to Gibbs structure

Point estimation

If entire history observed,**maximum a posterior (or maximum likelihood)** estimates for α, ϕ computable

Experiments

- ▶ Gibbs: parameter recovery
- ▶ Gibbs: scalability
- ▶ Point estimation with massive graphs

Parameter recovery

- ▶ Simulate 500 edges with fixed α
- ▶ Arrivals either $\mathcal{PYP}$ or Geom
- ▶ Observe final snapshot of the graph

Gen. arrival distn.	Inference model	$|\hat{\alpha} - \alpha^*|$	Pred. log-lik.
$\mathcal{PYP}(1.0, 0.75)$	$(\tau, \mathcal{PYP}(\theta, \tau))$	**0.046 ± 0.002**	**-2637.0 ± 0.1**
$\mathcal{PYP}(1.0, 0.75)$	$(\alpha, \mathsf{Geom}(\beta))$	0.049 ± 0.004	-2660.5 ± 0.7
Geom(0.25)	$(\tau, \mathcal{PYP}(\theta, \tau))$	0.086 ± 0.002	-2386.8 ± 0.1
Geom(0.25)	$(\alpha, \mathsf{Geom}(\beta))$	**0.043 ± 0.003**	**-2382.6 ± 0.2**

Scalability

▶ Simulate with fixed α and Geom(β) arrivals

	100 edges	10000 edges
$|\hat{\alpha} - \alpha^*|$	0.12 ± 0.01	0.01 ± 0.00
$|\hat{\beta} - \beta^*|$	0.02 ± 0.00	0.00 ± 0.00
Effective Sample Size	0.90 ± 0.04	0.75 ± 0.08
Runtime (s)	21 ± 0	2267 ± 2

▶ Runtime linear in #edges
▶ Most expensive Gibbs update is for **T**

MLEs for SNAP datasets

- SNAP datasets
- Fit point estimates for α, ϕ
- Fit: coupled $\mathcal{PYP}$, uncoupled $\mathcal{PYP}$ and Geom(β) arrivals

MLEs for SNAP datasets

Ask Ubuntu

- ▶ Estimates of $\mathcal{PYP}$ parameters vary significantly between coupled and uncoupled
  - ▷ $\hat{\theta}, \hat{\alpha} = 18080, 0.25$
  - ▷ $\hat{\theta}, \hat{\alpha}, \hat{\tau} = -0.99, -2.54, 0.99$
- ▶ Edge exchangeable models misspecified ($\eta > 2$)
- ▶ Using Geom estimates η well

Conclusion

- ▶ BNTL models are *flexible*
- ▶ BNTL models are *tractable*

References

[1] Jure Leskovec and Andrej Krevl. SNAP Datasets: Stanford large network dataset collection. `http://snap.stanford.edu/data`, June 2014.

[2] Diana Cai, Trevor Campbell, and Tamara Broderick. Edge-exchangeable graphs and sparsity. In *Advances in Neural Information Processing Systems*, pages 4249–4257, 2016.

[3] Harry Crane and Walter Dempsey. Edge exchangeable models for interaction networks. *Journal of the American Statistical Association*, 2017. In press.

[4] Benjamin Bloem-Reddy and Peter Orbanz. Preferential attachment and vertex arrival times. *arXiv preprint arXiv:1710.02159*, 2017.